

I ship AI features — then I *prove* they work.

LLM apps, agents, and workflow automation (RAG, LangGraph, n8n/Make) paired with the evaluation and regression harness that keeps them honest in production: Playwright, Pytest, k6, Promptfoo, DeepEval, LLM-as-judge. Most people build the AI or test it. I do both.

13+ yrs

QUALITY & AUTOMATION

3,759

TESTS · ONE PLATFORM

85

QUALITY RUNNERS BUILT

100×4

LIGHTHOUSE, OWN SITE

THE STORY THAT SAYS IT ALL

On 2026-08-15 my own quality platform refused its PROVEN verdict — 15 CVEs had drifted into production deps. The gate blocked *my* release. Patched and green the same afternoon; **both runs published verbatim**. A gate that never fires is decoration — mine fires, even on me.
agency.sageideas.dev/notes/gate-blocked-me.html

WHAT I BUILD (FIXED SCOPE, ~4 WEEKS, YOUR REPO)

- **LLM evaluation & AI quality** — golden sets, LLM-as-judge, safety battery (hallucination · injection · PII · toxicity), a CI gate that blocks bad prompt changes.
- **Test automation + CI** — risk-scoped Playwright/Pytest with release-manager evidence discipline; flake protocol that keeps red meaningful.
- **AI workflow automation** — n8n/Make/LangGraph builds with human approval points where the risk is, structured outputs, auditable logs.

WHY ME AND NOT A GENERALIST

CAPABILITY	AI BUILDER	QA ENGINEER	ME
Ships LLM features	✓	—	✓
Builds eval suites & judge gates	rare	rare	✓ shipped
CI regression discipline	—	✓	✓ 13 yrs
Certified (ISTQB CT-AI + TAE)	—	rare	✓ both

TRACK RECORD

The Home Depot (Fortune 50, 2012–21): Selenium framework for 2,300+ stores — regression 4h→75min, 99.5% stability. **HighStrike** (fintech, 2021–26): 500+ tests on live trading workflows, flake

rate 10%→<1%. **Sage Ideas** (2024–): the systems above, published with evidence.

Free 30-minute call. Bring the feature that scares you — leave with a plan either way.

agency.sageideas.dev